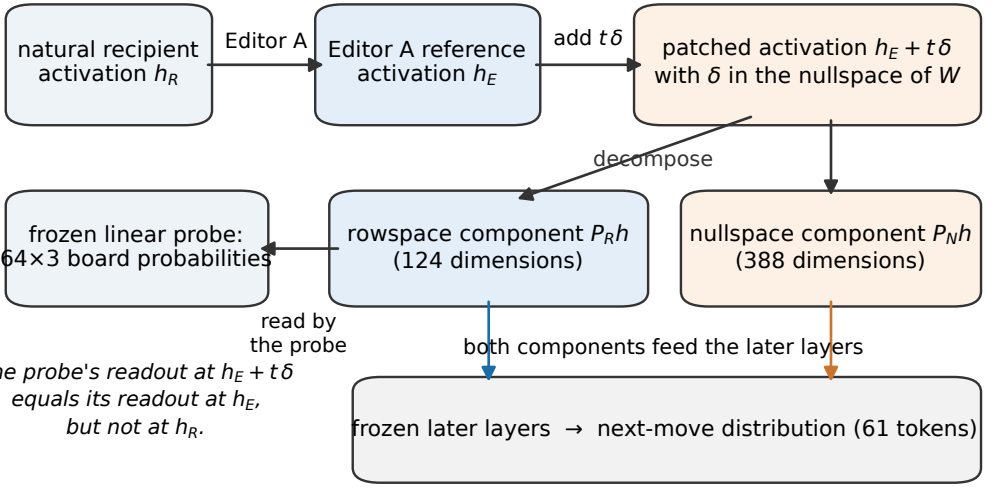
natural recipient activation h_R
Editor A
Editor A reference activation h_E
add $t\delta$
patched activation $h_E + t\delta$ with δ in the nullspace of W
decompose
frozen linear probe: 64×3 board probabilities
rowspace component $P_R h$ (124 dimensions)
nullspace component $P_N h$ (388 dimensions)
read by the probe
both components feed the later layers
The probe's readout at $h_E + t\delta$ equals its readout at h_E, but not at h_R.
frozen later layers → next-move distribution (61 tokens)